

DEVCONF.us

# Self Hosted LLMs

**Aakanksha Duggal**  
*Senior Data Scientist*

**Hema Veeradhi**  
*Principal Data Scientist*



**Hema Veeradhi**  
*Principal Data Scientist*



Emerging Tech, Office of  
the CTO



**Aakanksha Duggal**  
*Senior Data Scientist*



InstructLab



# Agenda

- Introduction to LLMs
- Open source LLMs
- Steps for building an LLM application
- Concept of self-hosting LLMs
- Setting up LLMs using Podman
- Demo
- Q&A



# Language Models

A **language model** is a type of machine learning model trained to conduct a probability distribution over words.




## Types of Language Models:

- **Statistical language models**
- **Neural language models**
  - **RNN**
  - **LSTM**
  - **Transformers**

# Large Language Models

- **Large Language Model - LLM** is just a larger version of a language model
- **WHY LLMs?**
  - **Quantitative** : Number of Parameters, **10-100 billion** parameters
  - **Qualitative** : Self-supervised learning




# LLM Applications

A Large Language Model is a type of Artificial Intelligence that is trained on a massive dataset of text and code. This allows the model to learn **statistical relationships between words and phrases which in turn allows it to generate text, translate languages, write creative content and answer your questions in an informative way.**

Here are common LLMs:

- **GPT3.5 and GPT 4**
- **Gemini**
- **Llama, Llama2**



# Open Source vs Closed Source Models

## List of Open Source Models available for commercial use

### Open Source models


- LLaMA-2 by Meta (claim to be)
- Falcon by Technology Innovation Institute in Abu Dhabi
- Mistral by Mistral AI

### Closed Source models

- GPT-4 by OpenAI
- Gemini by Google
- Claude by Anthropic


Project (maker, bases, URL)	Availability						Documentation				Access			
	Open code	LLM data	LLM weights	RL data	RL weights	License	Code	Architecture	Preprint	Paper	Modelcard	Datasheet	Package	API
BLOOMZ bigscience-workshop	✓	✓	✓	✓	~	~	✓	✓	✓	✓	✓	✓	✗	✓
LLM base: BLOOMZ, mT0      RL base: xP3														
Pythia-Chat-Base-7... togethercomputer	✓	✓	✓	✓	✗	✓	✓	✓	~	✗	~	~	✓	✗
LLM base: EleutherAI pythia      RL base: OIG														
•														
•														
•														
•														
LLaMA2 Chat Facebook Research	✗	✗	~	✗	~	✗	✗	~	~	✗	~	✗	✗	~
LLM base: LLaMA2      RL base: Meta, StackExchange, Anthropic														
Solar 70B Upstage AI	✗	✗	~	✗	~	✗	✗	✗	✗	✗	~	✗	✗	~
LLM base: LLaMA2      RL base: Orca-style, Alpaca-style														
Xwin-LM Xwin-LM	✗	✗	~	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	~
LLM base: LLaMA2      RL base: unknown														
ChatGPT OpenAI	✗	✗	✗	✗	✗	✗	✗	✗	~	✗	✗	✗	✗	✗
LLM base: GPT 3.5      RL base: Instruct-GPT														

# Levels of LLMs

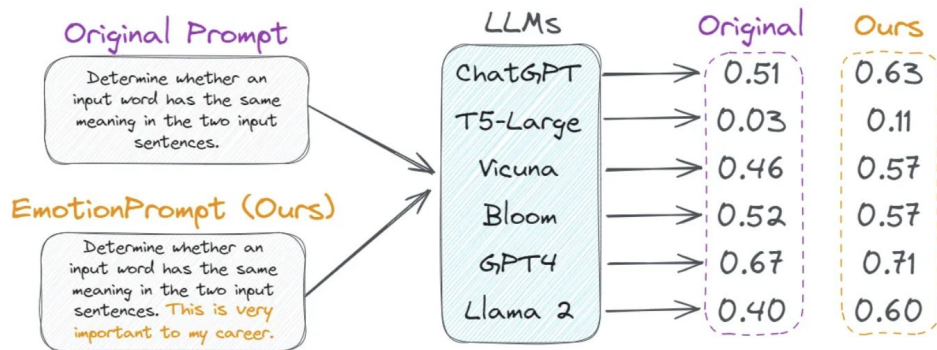






# Levels of LLMs



- “Let’s think step by step”
- “Take a deep breath and work on this problem step-by-step”
- “This is very important to my career”



# Levels of LLMs



## **Step 1:**

A Pre-trained LLM

## **Step 2:**

Update model parameters for a specific task

# Levels of LLMs

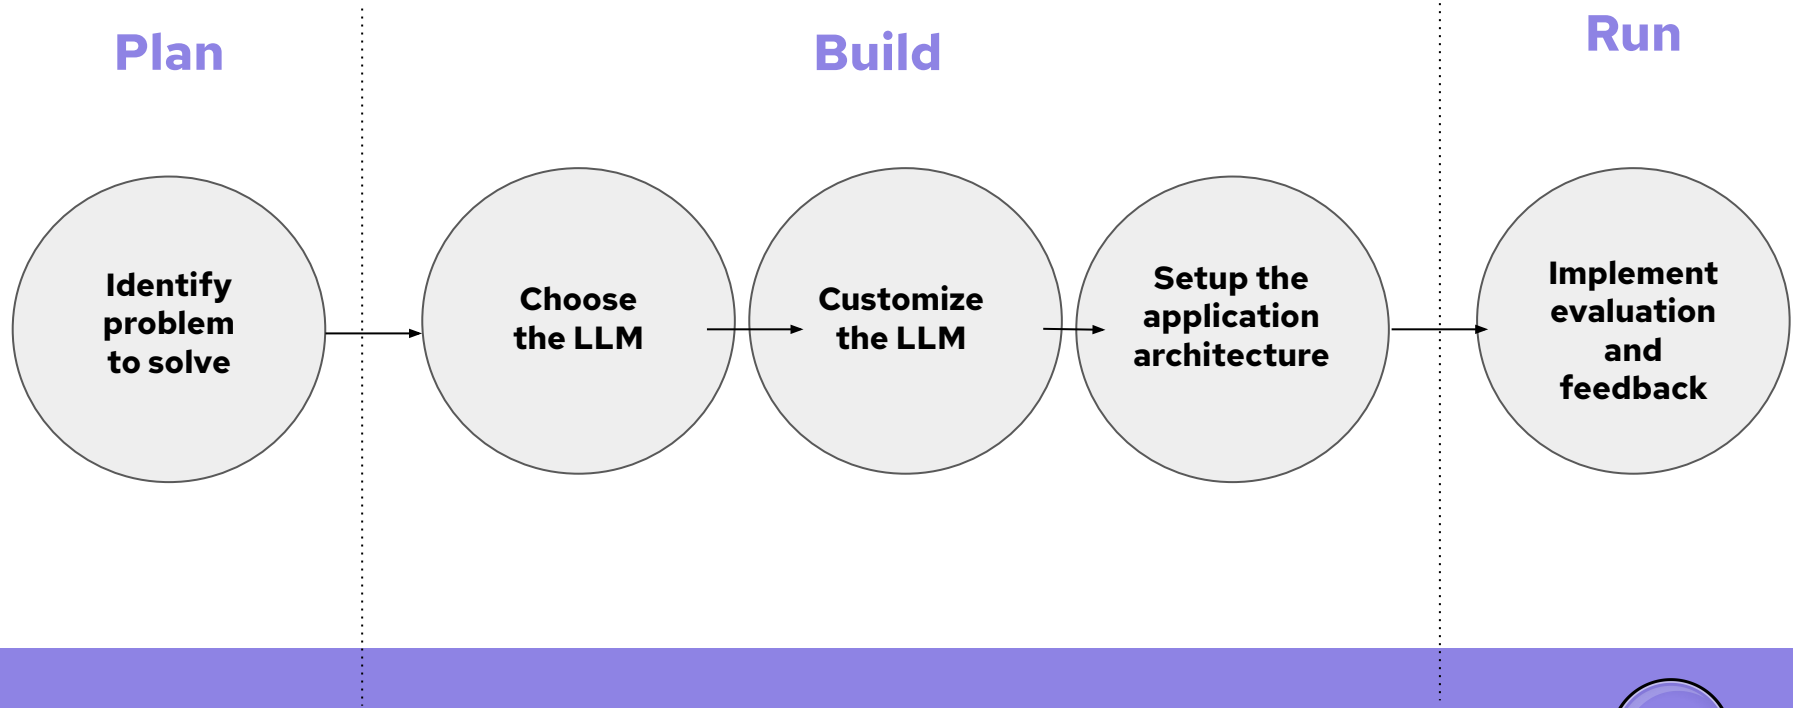
**Prompt  
Engineering**

**Model Fine  
Tuning**

**Build your  
own LLM**



# Steps for Building an LLM Application




# To self-host or not to?

- In the past year, the discussion surrounding LLMs has evolved, transitioning from **"Should we utilize LLMs?"** to **"Should we opt for a self-hosted solution or rely on a proprietary off-the-shelf alternative?"**
- Depending on your use-case, computational needs and engineering architecture availabilities you can decide whether to self-host your LLM
- Hosted models are **necessary** for privacy, reliability, or compliance



# Benefits of self-hosting LLMs

- Greater security, privacy, and compliance
- Customization
- Avoid vendor lock-in
- Save computational costs
- Easy to get started for those new to or just starting their journey with LLM




Schematic comparison of OpenAI GPT-3.5 and self-hosted LLMs

# Self-Hosting Containerized LLMs



**Hugging Face**

<https://huggingface.co/>



The image features an abstract background composed of several overlapping circles in various shades of purple, ranging from light lavender to deep indigo. The circles are arranged in a way that creates a sense of depth and movement. In the center of the composition, the word "DEMO" is written in a clean, white, sans-serif font. The text is positioned within one of the darker purple circles, which makes it stand out prominently. The overall aesthetic is modern and minimalist.



# Future Direction

- **Enhanced developer experience** enabling “non-data scientists” to follow a simple workflow for setting up and interacting with LLMs via microservices
- Implement a **seamless workflow** for transitioning from a local development environment to a production grade environment
- **End-end tooling**/framework for setting up LLMs locally for various applications such as text generation, document search, RAG applications etc

# Resources

- **GitHub repository:**  
<https://github.com/redhat-et/whisper-self-hosted-llm>
- **Slides:**  
<https://github.com/redhat-et/whisper-self-hosted-llm/tree/main/docs>
- **HuggingFace models:** <https://huggingface.co/ggerganov/whisper.cpp>
- **AI Lab recipes with Podman:**  
<https://github.com/containers/ai-lab-recipes>



Thank  
You!



[aduggal@redhat.com](mailto:aduggal@redhat.com)  
[hveeradh@redhat.com](mailto:hveeradh@redhat.com)